

Geographically-aware scaling for real-time persistent websocket applications.

Master's Project in Software Engineering

Lukasz Harezlak

lukasz.harezlak@gmail.com

Summer 2015, 41 pages

Supervisor: Tijs van der Storm
Host organisation: Instamrkt, <https://instamrkt.com>



UNIVERSITEIT VAN AMSTERDAM
FACULTEIT DER NATUURWETENSCHAPPEN, WISKUNDE EN INFORMATICA
MASTER SOFTWARE ENGINEERING
<http://www.software-engineering-amsterdam.nl>

Contents

Abstract	4
1 Introduction	5
1.1 Initial Study	5
1.2 Problem Statement	5
1.2.1 Research Questions	5
1.2.2 Solution Outline	5
1.2.3 Research Method	5
1.2.4 Hypothesis	6
1.3 Contributions	6
1.4 Related Work	6
1.5 Outline	7
2 Background	8
2.1 Scalability	8
2.1.1 Scalability Cost	8
2.1.2 Need for Scalability	9
2.1.3 Existing Approaches	9
2.1.4 Cloud Scalability	10
2.1.5 Data-layer Scalability	10
2.1.6 Websocket Scalability	12
2.1.7 Measuring Scalability	12
2.2 Geographical Distribution	12
2.3 Benchmarking	13
2.3.1 Benchmarking Crimes	14
2.4 Websocket Protocol	14
2.4.1 Client-Server Communication Improvements	14
3 The System Under Test	16
3.1 General Purpose of The System	16
3.1.1 Sample Use Case	16
3.1.2 Users of the System	16
3.2 Basic Architecture	16
3.3 Technology Stack	16
3.3.1 Programming languages	16
3.3.2 Database Technologies	17
3.4 System Scaleup Delay	17
3.5 Unique Aspects of The System	17
4 Experiment Outline	18
4.1 Goal of the experiment	18
4.2 Load Testing Framework	18
4.3 Kernel tuning	18
4.3.1 Measurements	18

4.3.2	Linux TCP stack	19
4.4	Baseline architecture on a local network	20
4.5	Cloud architecture setup	20
4.5.1	Route53	20
4.5.2	Autoscaling	20
4.6	Load balancing	21
4.7	Data Layer	21
4.7.1	Amazon Relational Database Service	21
4.7.2	Elastichache	21
4.7.3	Baseline architecture deployed in the cloud	21
4.7.4	Improved architecture deployed in the cloud	21
4.8	Experiment deliverables	21
5	Scalability Measurements	22
5.1	Selected metrics	22
5.1.1	EC2 Available metrics	23
5.2	Selected model	23
5.3	Baseline, Local network	23
5.3.1	Load Testing	24
5.4	Baseline, Deployed in the cloud	24
5.5	Improved, Deployed in the cloud	24
6	Experiment results	25
7	Evaluation	26
7.1	Statistical Significance	26
7.2	Threats to validity	26
8	Conclusion	27
9	Further work	28
10	Everything below that needs to be removed (except for bib)	29
11	Front Matter	30
11.1	Title	30
11.2	Author	30
11.3	Date	30
11.4	Host	30
11.5	Cover picture	31
11.6	Abstract	31
12	Core Chapters	32
12.1	Classic structure	32
12.2	Reporting on replications	33
12.3	L ^A T _E X details	34
12.3.1	Environments	34
12.4	Listings	34
13	Literature	37
13.1	Books	37
13.2	Journal papers	37
13.3	Conference papers	38
13.4	Theses	38
13.5	Technical reports	38
13.6	Wikipedia	39

13.7 Anything else	39
Bibliography	40

Abstract

This section summarises the content of the thesis for potential readers who do not have time to read it whole, or for those undecided whether to read it at all. Sum up the following aspects:

- relevance and motivation for the research
- research question(s) and a brief description of the research method
- results, contributions and conclusions

Kent Beck [[JBB⁺93](#)] proposes to have four sentences in a good abstract:

1. The first states the problem.
2. The second states why the problem is a problem.
3. The third is the startling sentence.
4. The fourth states the implication of the startling sentence.

Chapter 1

Introduction

Software scalability. Cloud scalability. Why is it important Importance of measuring in the clouds and testing optimal architectures.

1.1 Initial Study

What we found out researching.

1.2 Problem Statement

what's the best architecture that answers that best: how to deliver data to geographically distributed users with minimal, consistent and manageable latency how to react scales up and down in response to demand changes (WEBSOCKETS - DISCONNECT SESSIONS? WHEN? This provides some additional challenges when it comes to websockets, since in order to stop an instance one needs to make sure it does not serve any sessions. [GB14]) Moreover, they do not consider the cost and degree of utilization of the employed resources within a data center. [GB14] according to the models based on the set of metrics described in detail later ?? Also, there is an inherent limitation for scaling a websocket application - a number of TCP ports (and file descriptors available to a ws/wss connections) on a server instance. Hardware limitations are different in an http-based application.

1.2.1 Research Questions

1. Does architecture with geographically aware scaling case can produce better results than the baseline architecture?
2. What is the preferred architecture decomposition for a system with given characteristics - stateless, deployed in the cloud, dynamically scaled up and down, with persistent connections, clients distributed globally and uncacheable data generated in real time?

1.2.2 Solution Outline

How our solution works in short.

1.2.3 Research Method

Software engineering is a relatively difficult field to investigate in terms of lack of clarity on how to do it. Most of the problems are of design, rather than pure scientific, nature. West Churchman coined a specific term for these kind of modern problems - wicked (since they are resistant to resolutions) [MMPROJ1]. Eastbrook et al claim it is often difficult to identify the true underlying nature of the research problem and thus the best method to research it [MMPROJ2]. In their work, they name and compare five most classes of research methods to select from: controlled experiments,

case studies, survey research, ethnographies, action research. They help to select the method by first establishing the type of research question being asked - existence question (Does X exist?), description and classification question (What is X like?), and descriptive-comparative questions (How does X offer from Y?). The questions of this research are of the last type. The authors suggest pinpointing upfront what will be accepted as a valid answer to the research question [MMPROJ2]. The detailed description of each of the methods helps me to settle for the controlled experiment. Its well-suited for testing a hypothesis where manipulating independent variables has an effect of dependent ones, which is exactly the case of me research. The manipulated variable is architecture decomposition, and the measured ones are determined by scalability measurement models described below.

Research Difficulty

The experiments will be performed in a shared cloud environment, which is inherently unpredictable and changing. Designing and executing a test yielding statistically relevant results where all the necessary variables are controlled (within reasonable boundaries) will be a challenge. The focus of the project is on scaling a websocket application. This is a relatively new technology, thus finding proper scientific coverage is not trivial. Some of the necessary development and data collection might be in relatively low-level technology. The geo-location with reasonable accuracy of the clients might prove difficult too. Overall, its a complex project requiring knowledge of both hardware (protocol), software (architecture) and consisting of multiple disciplines. even linking order in the compiler can change the performance!!! [one of CRIMES]

1.2.4 Hypothesis

The correct geographical decomposition of application stack can lead to vastly improved performance in comparison with baseline architecure(TODO:S????).

1.3 Contributions

Case study on AWS. Architecture decomposition analysis. Deliverable piece of code - metrics framework and auto-scaling scripts.

1.4 Related Work

The big paper on measuring scalability in the clouds. custom routing between regions. A lot on page 14/21 MMPROJ. The big paper on deploying different scalability structures [GB14]. Software engineers need to design for the cloud, not only deploy in the cloud. This is even more important when using multiple data centers situated in different legislative domains; constructed with different hardware, network, and software components; and prone to different environmental risks. Both Route 53 and ELB do not consider applications regulatory requirements when selecting a data center site. Moreover, they do not consider the cost and degree of utilization of the employed resources within a data center. because they want custom metrics routing they Selecting the cloud: As a first step, the user authenticates to one of the entry points. At this point, the entry point has the users identity and geographical location (extracted from the IP address). As a second step, the entry point broadcasts the users identifier to the admission controllers of all data centers. We call this step matchmaking broadcast. As a first step, the user authenticates to one of the entry points. At this point, the entry point has the users identity and geographical location (extracted from the IP address). As a second step, the entry point broadcasts the users identifier to the admission controllers of all data centers. We call this step matchmaking broadcast. The admission controllers respond to the entry point whether the users data are present and if they are allowed (in terms of legislation and regulations) to serve the user. In the response, they also include information about costs within the data center. Based on the admission controllers responses, the entry point selects the data center to serve the user and redirects him or her to the load balancer deployed within it. The entry point filters all clouds that have the users data and are eligible to serve him or her. If there is more than one such cloud, the entry

point selects the most suitable with respect to network latency and pricing. If no cloud meets the data location and legislative requirements, the user is denied service. After a data center is selected, the user is served by the AS and DB servers within the chosen cloud as prescribed by the standard three-tier architecture. they do: sticky load balancing it is not beneficial to terminate a running VM ahead of its next billing time. It is better to keep it running until its billing time in order to reuse it if resources are needed again. their objectives - cost minimization and latency minimization they use front VMs and each cloud runs an ELB an Admission Controller latency approximation GeoLite db (ip-geo coordinates) compute latency using PingER (To approximate the latency between a user and a cloud, we select the three pairs of PingER hosts that are closest to the user and the cloud, respectively, and define the latency as a weighted sum of the three latencies between the hosts in these three pairs) CloudSim Small mentions of the others.

1.5 Outline

Here we outline the structure of the thesis. A short paragraph on what happens in each chapter.

Chapter 2

Background

2.1 Scalability

Scalability seems to be a notion that everyone intuitively grasps, but has difficulties when it comes to clear explanations. In my literature study I came across a few definitions, which might help with that, of which three can be found below:

[WS04] Scalability is a measure of an application systems ability to, without modification, cost-effectively provide increased throughput, reduced response time and/or support more users when hardware resources are added.

[WG06] Scalability is an ability of a system to handle increased workload (without adding resources).

In light of this definition we would talk about a scalability failure if one of the following occurred:

- Address space was exceeded,
- Memory was overloaded,
- Available network bandwidth was exceeded,
- etc.

[WG06] Scalability is an ability of a system to handle increased workload by repeatedly applying a cost-effective strategy for extending a systems capacity.

According to this definition, one could determine system scalability failure if a given resource got overloaded or exhausted an adding capacity to this resource would not result in a proportional ability to handle additional demand, e.g.:

- an additional processor will not contribute to meeting the higher demand if handing of that processor entails an overhead).
- a newly added server instance might not contribute to handling a higher user demand if slows down the routing process.

2.1.1 Scalability Cost

Scalability is generally desired in the software systems, yet, as all architectural software decisions come at a cost [Hei10], one must be aware of the trade-offs usually associated with it. Weinstock and Goodenough [WG06] point these out:

- performance and scalability (non-scalable system will often demonstrate degrading performance with increasing demand, but scalable systems require performance sacrifice on lower usage levels),

- cost and scalability (designing a system to be scalable up-front entails additional costs),
- operability and scalability (it is difficult for humans to operate large systems),
- usability and scalability (it may be possible to increase servable demand with limiting system's service scope - e.g. removing personalization and displaying generic, cacheable data),
- data consistency and scalability (higher scalability can be achieved if system allows for data inconsistencies).

2.1.2 Need for Scalability

As the globalization and internetization progresses, more and more systems are expected to be capable of serving millions of globally distributed users. A single server instance often cannot live up to that task and thus application needs to be divided and distributed in multiple smaller chunks. This division happens on different application layers, and different parts of the system have to communicate and synchronize with each other.

In case of many software systems (the system under test 3 included), user traffic, and with it the need for system services and resources, varies significantly. A need to be able to scale up and down dynamically in response to traffic arises from this.

Huge parts of the internet are shifting towards real-time. This trend is giving rise to new technologies for exchanging messages between clients and servers in the client-server architecture. Traditionally, client would send a request to a server and receive a response. This is hugely inefficient when there is a need for continuous bidirectional exchange of messages. As an improvement, new mechanisms for server-client communication have been introduced recently 2.4.1 to enhance the process and reduce overhead. One of them - websockets - are a huge next step on this path, but introduce new challenges. One of them is scalability of applications which make use of this technology.

2.1.3 Existing Approaches

There exist multiple scalability vectors for web applications. Different decompositions of application stack can be applied to achieve the goal of scalability. A few traditional approaches of tackling an issue like that exist already [Lei08]:

Scaling up. Increasing volume of system allocated resources. As internet is unreliable and so called "middle-mile bottlenecks" exist, a web application end-user latency and throughput experience is not fully deterministic. Traffic levels fluctuate tremendously, so the need to provision for peak traffic levels means that expensive infrastructure will sit underutilized most of the time. It is easy to implement, but costly, even extremely when you start pushing at current hardware limits.[Qve10]

Scaling out. Scaling out horizontally - increasing the number of units of resources comprising the system. Cost of hardware can be reduced dramatically this way. In a web application, one can deploy multiple instances of servers. Different types of balancing can be applied to distribute traffic among them: application layer balancing, business load balancing, and anticipating load. The overhead of parsing requests in the application layer is high thus limiting scalability compared to load balancing in the transport layer. Client state needs to be stored in a layer shared between the web servers.[Qve10]

Content Delivery Networks. These only handle static assets. Communication in our system 3 happens mainly over websockets, which are not supported by CDNs. CDNs can be divided into Big Data Center CDNs and Highly Distributed CDNs, which put application data within end-user ISPs.[Lei08]

Peer to peer networks. An architecture different from client - server, where users communicate with each other directly. It handles adding and removing nodes to and from the network dynamically very well.

2.1.4 Cloud Scalability

Cloud computing has become virtually ubiquitous. All biggest internet services are either deployed to a cloud, or run their proprietary cloud systems. A lot of companies running proprietary clouds also make them available for hosting to the external clients. Netflix¹ and Spotify² (both deployed to Amazon's AWS) are examples of the first approach, with Microsoft(Azure)³, Google(Google Cloud Platform)⁴ and Amazon(Amazon Web Services)⁵ being the example of the second. General cloud capabilities and the cloud stack I am working with for the scope of this project is described in a dedicated chapter 4.5.

One of the biggest advantages of deploying one's architecture to the cloud that aforementioned companies offer is the ability to dynamically scale up and in response to application traffic changes.

As Grozev and Buyya [GB14] put it, to fully facilitate cloud capabilities, software engineers need to design for the cloud, not only to deploy to it.

2.1.5 Data-layer Scalability

Data layer scalability is an important part of system scalability. In a distributed system, multiple system agents share the data. There exist multiple strategies for operating on a shared, distributed data layer.

Data layer is often a performance bottleneck because of requirements for transactional access and atomicity - it is hard to scale out when system uses a relational data store [GB14].

CAP Theorem put forward by Eric Brewer [GL02] states that it is impossible for a distributed system to provide all of the following guarantees:

- consistency,
- availability,
- partition tolerance.

Therefore, a trade-off needs to be made, depending on stakeholder priorities, which of the three to give up.

ACID stands for Atomicity, Consistency, Isolation, Durability and is a set of properties guaranteeing reliable processing of database transactions. It has been a guideline in designing multiple database systems. The term was originally coined by Haerder and Reuter in 1983 [HR83].

BASE stands for Basically Available, Soft State, Eventually Consistent [LD12]. It's a complement of ACID. Author claims we lack precise metrics to measure BASE aspects and that's why every system implements eventual consistency differently.

Most importantly for this research, the author [LD12] explains *MySQL Cluster* - it performs much like an ACID database but with the performance benefits of a cluster. Aside from that, *MySQL Replication* can be put to use. It can be configured at multiple topologies, not only the basic master-slave; consistency differs per configuration.

Many solutions and strategies for dealing with database scalability have emerged, including *NoSQL* and *NewSQL* databases, data replication and sharding [Amz02].

Cooper et al touch on that subject in their work [CSE⁺10]. They claim that in scaling out the database layer one should aim for elasticity (dynamically adding capacity to a running system) and high availability. These are hard to achieve using traditional database systems. They show that new protocols are being developed to address that issue, such as that: *two-phase commit protocol* (provides atomicity for distributed transactions) and *paxos*.

¹<https://aws.amazon.com/solutions/case-studies/netflix/>

²<https://aws.amazon.com/solutions/case-studies/spotify/>

³<https://azure.microsoft.com/>

⁴<https://cloud.google.com/>

⁵<https://aws.amazon.com/>

The authors [CSE⁺10] also give an overview of different database systems, including *PNUTS*, *BigTable*, *HBase*, *Cassandra*, *Sharded MySQL*, *Azure*, *CouchDB*, and *SimpleDB*.

In their work [CSE⁺10], we can find enumeration of classic data-related scalability trade-offs:

- read performance and write performance,
- latency and durability,
- synchronous and asynchronous replication,
- data partitioning (column and row-based storage).

The technology stack I have been working with in the scope of this project consists of *MySQL* as persistent storage and *Redis* as a key-value cache. It is described in details in a dedicated chapter 3.3.

MySQL Scalability

MySQL White paper [Ora15] gives us a good overview of how scalability works in case of MySQL Cluster. Authors suggest starting the design of the scaling process by identifying which characteristic the application possesses: lots of write operations, real-time user experience, 24x7 user experience or agility and ease-of-use. The System Under Test 3 falls into the first and second categories.

They claim to support (among others) auto-sharding for write-scalability, active / active geographic replication and online scaling and schema upgrades. Geographic replication offers distribution of clusters across remote data centers, which they claim helps reduce latency (which is extremely important case of The System Under Test 3).

Authors show how MySQL Cluster is optimized for real-timeness:

- data structures are optimized for in-memory access,
- persistence of updates runs in the background,
- all indexed columns are stored in memory.

According to the white paper [Ora15], MySQL Cluster is very well-suited for on-line, on-demand scaling.

MySQL can also be used to scale in unconventional ways. Ruffin et al, in Social-Data Storage-Systems [RBR11], mention that Twitter uses MySQL as a key value store.

They show how MySQL can be scaled horizontally by both sharding and data replication. They also indicate that the more structured the RDBMS data, the harder it is to scale horizontally. MySQL is optimized for writes, since only one record in one table is touched, whereas reads can prove expensive if they contain joins, especially spreading across multiple cluster nodes. Facebook and Twitter solved it by putting a cache on top of MySQL [RBR11].

Redis Scalability

Single Redis installation is said to be of limited scalability, because of the fact that for good performance the whole data set should fit into memory [RBR11]. Redis offers cluster⁶ and replication⁷ solutions.

In a Redis Cluster, data is automatically sharded (not replicated) among multiple nodes. One has control over sharding; it is possible to ensure that certain data points end up on the same (or on a given) node. Redis Cluster does not guarantee strong consistency - under certain conditions the Cluster can lose writes that were acknowledged to the client. Support for synchronous writes exists, through the command *WAIT*, which highly (but not entirely) decreases the likelihood of lost writes. Nevertheless, using it is discouraged unless absolutely necessary.

Redis can also be scaled for reads using a simple replication in a single master - multiple slaves topology.

⁶<http://redis.io/topics/cluster-spec>

⁷<http://redis.io/topics/replication>

2.1.6 Websocket Scalability

An introduction into push-base communication over the internet can be found in Agarwals work - Toward a Push-Scalable Global Internet [Aga11]. The key message in the article is that push message delivery on the World Wide Web is not scalable for servers, intermediate network elements, and battery-operated mobile device clients. And yet, most of modern day websites have highly dynamic content updated even up to multiple times a minute (very much so for The System Under Test 3).

Most of internet communications happens over HTTP Running over TCP. Real-time message delivery requires an always-on connection from the server to the client. HTTP proxies have limited memory and tcp ports, are shared among multiple users. Servers need to be provisioned in order to maintain active tcp connections from large populations of user clients. These all provide challenges that need to be dealt with in scalable applications [Aga11].

On a more positive note, Cassetti and Luz [CL12] claim that overhead introduced by the websocket protocol and websocket API is rather small as compared to other communication methods. Furthermore, one data intensive applications can achieve superior bandwidth and performance when using websockets.

My own research has proven preparing websocket-based communication scalability to be difficult and tedious. Most of the cloud scalability stack that is available and that I have been working with was designed and built to support http-based – and not websocket-based – communication. On top of that, one needs to tweak the host operating system kernel in multiple ways to increase the number of concurrent websocket connections the system can maintain. Details of the tuning process are described later 4.3.

The Websocket Protocol is described in details in a dedicated section 2.4.

2.1.7 Measuring Scalability

Measuring scalability proves a challenge, since there is no single physical quantity or unit which the community has accepted as a scalability measure.

Measuring and researching in a cloud environment is difficult, because, by definition, the researcher has no control over the hardware that his system is running on. Clouds are also inherently non-deterministic. Nevertheless, as Sobel et al. describe in their work [SSS⁺08], even in cloud computing environments, where researchers have little control over network topology or hardware platform, understanding the performance bottlenecks and scalability limitations imposed by the offered infrastructure is valuable.

As this is a crucial topic to my research, this topic is explored in detail in a dedicated chapter 5.

2.2 Geographical Distribution

With the rise of global internet services, single applications have to server users who are distributed around the whole globe. With the distribution of the users, comes distribution of the data. Multiple solutions exist that, through data distribution, aim at improving the *Quality of Service*.

The topic is nicely introduced by Tom Leighton in his article *Improving performance of the internet* [Lei08]. He provides a few arguments why it is important to keep data as close to end users as possible. Apart from obvious latency benefits, by doing this, one reduces the chances of suffering from a big *middle mile* provider outage. He introduces a scale to reason about internet locality.

Scale Name	Range	Latency Range	Typical Packet Loss
local	less than 100 miles	1ms	0.6%
regional	500 - 1000 miles	15ms	0.7%
cross-continent	3000 miles	50ms	1.0%
multi-continent	6000 miles	100ms	1.4%

Table 2.1: Internet Locality

As the table 2.1 shows, the longer data must travel through the middle mile, the more it is subject to congestion, packet loss, and poor performance [Lei08]. This is why companies try to locate servers close to end users (on a scale that cannot be reproduced in this research - requires finer control over infrastructure). The lowest granular scale I am able to work with within this research is *regional*. Leighton writes that big websites have at least two geographically dispersed mirror locations to improve performance, reliability and scalability.

In his work, he includes a set of guidelines to consider when performing geographical scalability:

- reduce transport layer overhead (I am achieving this with usage of the websocket protocol),
- prefetch embedded content,
- assemble pages at the edge (even on an end client machine),
- offload computations to the edge,
- ensure significant redundancy in all systems to facilitate failover,
- use software logic to provide message reliability.

Ed Howorka in his interesting paper *Colocation beats the speed of light* [How15] about trading system geographical distributions focuses on the best placement of servers for traders trading on multiple exchanges. His paper demonstrates that traders gain nothing by positioning their computer at the midpoint between two financial exchanges. He claims that every algorithm on a central machine talking to surrounding servers (users in case of the System Under Test 3) can be replaced by colocated servers (located in geographical proximity to users). Furthermore, he implies and sets out to prove that server colocation provides a better solution for high-speed applications (as opposed to using a big, centralized server located in the middle).

2.3 Benchmarking

benchmarking run a few times for caches to fill up etc

[CSE⁺10] important to have random distributions(uniform, zipfian, latest, multinomial)

[MMPROJ2] <https://homes.cs.washington.edu/~bornholt/post/performance-evaluation.html> computer science experiments tend to be among the most non-deterministic and prone to omitted-variable bias Computer systems have many moving parts in both hardware and software, researchers are unlikely to have a complete understanding of every part in those systems, and controlling each of those components is often next to impossible. linking order can significantly bias the results of a benchmark the size of your UNIX environment variables biases performance - we made sure this are consistent across tests

Most researchers Ive met tend to rebuild their infrastructure over and over again, writing new scripts to parse, tabulate and plot their results for every experiment. This approach clearly doesnt scale. This is where we can come in and contribute to the scientific and open source communities. Plan to release the deployment and whatever is possible.

[MMPROJ5] http://www.sosy-lab.org/dbeyer/Publications/2015-SPIN.Benchmarking_and_Resource_Measurement.pdf benchmarking good and well established for very specific tasks (<https://www.spec.org/> 2 <http://www.tpc.org/> 3 <http://nlrp.ipd.kit.edu/>), our case more complex (as describd below) non-deterministic aspects (networks, multi-threading etc.) Results are reproducible if it is guaranteed that the same results can be obtained later again (assuming a deterministic tool) by re-running the benchmarks on a machine with the same hardware and the same software versions. Reproducibility of experimental results requires reliable measurement. We call a measurement reliable, if the measurement method ensures accuracy (small systematic and random measurement error, i.e., no bias or volatile effects, resp.) and sufficient precision

Benchmarking a distributed tool (3) is much more complex and out of scope. [3rd paper from http://www.sosy-lab.org/dbeyer/Publications/2015-SPIN.Benchmarking_and_Resource_Measurement.pdf] extremely difficult to get all variables under control in IT - hardware, all the layers of software, non deterministic stuff like networks, etc. Don't even get me started on clouds where you basically give up control over hardware to the cloud host.

2.3.1 Benchmarking Crimes

[Hei10] Selective benchmarking not covering the full evaluation space not evaluating potential performance degradation Progressive criterion - performance actually does improve significantly in area of interest Conservative criterion - performance does not significantly degrade elsewhere both need to be demonstrated - we wrote our testing scenarios to cover the full area of interest Reality is that techniques that improve performance generally require some degree of extra work: extra bookkeeping, caching, etc. These things always have a cost, and it is dishonest to pretend otherwise. This is really at the heart of systems: it's all about picking the right trade-offs. A new technique will therefore almost always introduce some overheads, and you need to demonstrate that they are acceptable. - how do we fit in here?

benchmark subsetting without strong justification - WARNING SIGNS: we picked a representative subset, typical results are shown (reads as cherry picked to fit our results), I want to understand why you picked the particular subset [of benchmarks]. If your selection contains many obscure or outdated devices, or is heavily biased towards X and Y drivers, then I suspect that you have something to hide. Selective data set hiding deficiencies. Luckily we have statistical tools to detect that (chess data skeptic podcast)

Same dataset for calibration and validation Calibrate the system first (using calibration workload). Use evaluation workload to show how accurate the model is. Both workloads need to be totally disjoint. - TODO: DOES THIS APPLY TO OUR CASE?

No indication of significance of data. Raw averages, without any indication of variance. At least standard deviations must be quoted. If in doubt use student's t-test to check significance.

Benchmarking of simplified simulated system - WE'RE DOING THE REAL DEAL.

Inappropriate and misleading benchmarks, Unfair benchmarking of competitors - We care about our system's best setup so no incentive to lie.

Relative numbers only, Reads as I am covering that the results are really bad or irrelevant. - happy to provide absolutes but sometimes the unit will be made up so numbers only make sense in context

No proper baseline - Often the state-of-the-art solution, we couldn't find data to prove that what we are comparing against is bad. Occam's razor, simple solution proves best - only known limitation is the number of sockets.

Arithmetic mean for averaging across benchmark scores The proper way to average (i.e. arrive at a single figure of merit) is to use the geometric mean of the normalised scores [http://dl.acm.org/citation.cfm?id=5673]

[MMPROJ3] http://dl.acm.org/citation.cfm?id=5673 Using the arithmetic mean to summarize normalized benchmark results leads to mistaken conclusions that can be avoided by using the preferred method: the geometric mean.

2.4 Websocket Protocol

2.4.1 Client-Server Communication Improvements

M. Franklin and S. Zdonik provide an insight into the history of push-based technologies. Their paper from 1998 [MMPROJ12] classifies communication mechanisms into aperiodic pull, periodic pull, aperiodic push and periodic push.

previous: ajax, long polling, short polling, websocket [Aga11] Before the rise of the websocket protocol, different techniques were used to improve communication between server to client. Among them, the following can be listed (some of them serving different purposes): ajax - a request / response model; was an improvement since the page didn't have to be refreshed anymore to get new data, short polling - using a timer to regulate sending requests to server, similar to refreshing the page, useful when data doesn't change too often, long polling - now considered to be a workaround of preventing creating connections for each request, keeps a connection artificially alive for some time, clients have to reconnect periodically, webRTC - a peer to peer connection server-side events - only allow for a server-initiated communication.

websocket draft - RFC6455

[http://stackoverflow.com/questions/4852702/do-html-websockets-maintain-an-open-connection-for-each-client-does-this-scale/25340220#25340220] Each TCP connection in itself consumes very little

in terms server resources. Often setting up the connection can be expensive but maintaining an idle connection it is almost free. The first limitation that is usually encountered is the maximum number of file descriptors (sockets consume file descriptors) that can be open simultaneously. This often defaults to 1024 but can easily be configured higher. http connections more expensive: - Each HTTP connection carries a lot of baggage that isn't used most of the time: cookies, content type, content length, user-agent, server id, date, last-modified, etc. Once a WebSockets connection is established, only the data required by the application needs to be sent back and forth. - TYPICALLY: HTTP servers are configured to log the start and completion of every HTTP request taking up disk and CPU time. It will become standard to log the start and completion of WebSockets data, but while the WebSockets connection doing duplex transfer there won't be any additional logging overhead (except by the application/service if it is designed to do so) - interactive applications that use AJAX either continuously poll or use some sort of long-poll mechanism. WebSockets is a much cleaner (and lower resource) way of doing a more event'd model where the server and client notify each other when they have something to report over the existing connection. obviously don't support cdn

As for decreased latency via WS vs. HTTP, it's true since there's no more parsing of HTTP headers beyond the initial WS handshake. Plus, as more and more packets are successfully sent, the TCP congestion window widens, effectively reducing the RTT.

Think of it this way: what is cheaper, keeping an open connection, or opening a new connection for every request (with the negotiation overhead of doing so, remember it's TCP.)

<http://serverfault.com/questions/48717/practical-maximum-open-file-descriptors-ulimit-n-for-a-high-volume-system> the number of client connections that a server can support has nothing to do with ports in this scenario, since the server is [typically] only listening for WS/WSS connections on one single port. I think what the other commenters meant to refer to were file descriptors. You can set the maximum number of file descriptors quite high, but then you have to watch out for socket buffer sizes adding up for each open TCP/IP socket. **MAKE SURE THIS IS CONSISTENT THROUGHOUT THE PAPER.** If the file descriptors are tcp sockets, etc, then you risk using up a large amount of memory for the socket buffers and other kernel objects; this memory is not going to be swappable.

Chapter 3

The System Under Test

real time prediction parmet parimutuel pools directed contracts sort of a stock market for real-time prediction on anything, e.g. live sport

3.1 General Purpose of The System

first applicatoin - live sports low manageable latencies important

3.1.1 Sample Use Case

A good example e.g. can be predicting the outcomes of certain drives in the football match as they happen. A sample case: Manchester United - Liverpool live game, live data coming from UK servers, introduced into the system through a UK node, most users (who also generate data that needs to be distributed) in China (10k), India(10k), Australia(3k). Where should websocket servers which distribute messages spun up for minimal latencies for all clients? Is it faster to spin up db replicas on that nodes too?

3.1.2 Users of the System

in stadium on couch dekstop tablet mobile (The bandwidth of a wireless network is more physically constrained in size and scope than a wired network. With less overall bandwidth to work with, the number and size of connections that are using that bandwidth must be managed more carefully because they have a greater impact on performance. This is why latency is more of an issue or wireless networks than on wired ones. <http://developer.att.com/application-resource-optimizer/docs/best-practices/multiple-simultaneous-tcp-connections>) Users of the platform are mostly on mobile networks (that often drop), so reconnecting them quick to the right (providing lowest latencies) instance is important. lumpy demand - it comes in 2-hour-long spikes and then can go quiet for days

3.2 Basic Architecture

Diagram

3.3 Technology Stack

Python + redis + mysql + redis (cluster) all versions cloud stack described here: [4.5](#)

3.3.1 Programming languages

python

3.3.2 Database Technologies

Mysql Cluster

write scalable, important for us, multi master, automatic / manual sharding distributed, shared-nothing data that provides very high, scalable performance for applications that largely use primary key access and have high concurrency, but it incurs a cost in network access when accessing data between a MySQL Server and tables distributed across Data Nodes In terms of CAP, MySQL Cluster will sacrifice availability to maintain consistency in the event of an unsurvivable network partition. at least 3 machines (otherwise split brain problem) data nodes (ndbd) - need memory, single-threaded sql nodes (mysqld) - need cpu, multithreaded Important: Will MySQL Cluster out of the box run an application faster than my existing MySQL database? The reality is probably not without some optimizations to the application and the database. However, by following the guidelines and best practices in the Guide to Optimizing Performance of the MySQL Cluster Database1 , significant performance gains can be realized, especially as you start to scale write operations while maintaining very high levels of availability and low latency. [<http://openquery.com.au/files/mysql-cluster-intro-use.pdf>]

3.4 System Scaleup Delay

Delays: up to 2 minutes to get metrics, up to ????? minutes for instance to be accessible to clients + ???? for DNS changes to propagate (TTL set to 60 seconds but TODO VERIFYF) Unfortunately there are a large number of (misbehaving) DNS servers out there that dont properly obey TTLs on records and will still serve up stale records for an indefinite amount of time. [<http://engineering.chartbeat.com/2014/01/02/part-1-lessons-learned-tuning-tcp-and-nginx-in-ec2/>]

3.5 Unique Aspects of The System

critical data flows over websockets Users receiving data shared locally should receive it at the same time as data shared globally (with as low a latency as possible). This creates a need for globally consistent and manageable latency between end user and the system. Connected users and sources of data are geographically changing. Users geographical center of mass is changing for each peak of demand. Extremely lumpy demand (peaks lasting around 2 hours). This creates a need for being able to quickly scale up and scale down. New data is generated every few seconds by the users so caching the content and distributing geographically is difficult.

Chapter 4

Experiment Outline

1. Baseline architecture on a local network.
2. Baseline architecture deployed in the cloud.
3. Improved architecture in the cloud.

All of them measured with the same set of metrics with a prepared framework for gathering them. Details below.

4.1 Goal of the experiment

The goal of the project is to design a scalability framework for a real-time persistent websocket distributed application (the system). The core researched topic will be whether a systems awareness of clients geographical distribution can improve the system performance according to selected metrics, in comparison with traditional approaches.

The goal of the project is to see if the proposed architecture decomposition can perform better (quantitatively, according to the selected metric model) than the baseline architecture in serving geographically dispersed clients. Approaches used to scale simple http applications cannot always be translated to websocket applications since the communication protocol differs. Websockets put a different kind of strain on the server machines since these need to keep the connection opened on a port for a prolonged period of time rather than simply open, server and close (as is the case with http). Along the way, an answer needs to be found what level of decoupling provides best performance on each layer of a stateless persistent system. One of the properties of a system of that kind is that key value stores come under heavy load since this is where the state resides. A good solution for distributing (sharding / replicating) these also needs to be found. Same goes for persistent storage. TODO: REFERENCE CURRENT APPROACHES HERE.

4.2 Load Testing Framework

all the analyzed tools, list also in lab notes from may 11 [FOR EACH WHY WASN'T SATISFACTORY] jmeter, thor (low extensibility which we needed), autobahn (max available connections but no further control), gatling-websocket, <http://www.opensourcetesting.org/performance.php> tsung - looked promising, highly scalable erlang, but hard to understand what was going on, community was not active enough, wsbench

4.3 Kernel tuning

4.3.1 Measurements

etsy, logster, graphite etc.

4.3.2 Linux TCP stack

all lab notes from 08.05 all websocket updates <http://serverfault.com/questions/48717/practical-maximum-open-file-descriptors-ulimit-n-for-a-high-volume-system> Also, and very important, you may need to check if your application has a memory/file descriptor leak. Use lsof to see all it has open to see if they are valid or not. Don't try to change your system to work around applications bugs.

Although the causes for such symptoms can vary, there's one scenario that can cause a complete lock of systems handling a very large number of web requests per second without any hint of what's going on: TCP/IP port exhaustion. [<https://www.outsystems.com/forums/discussion/6956/how-to-tune-the-tcp-ip-stack-for-high-volume-of-web-requests/>]

describe close_wait, time_wait etc

only 65k per (tcp 16 bit) ip address, one can add virtual ips to machines and up the number this way (recent lab notes)

increasing the number of clients The solution is more quadruplets⁵. This can be done in several ways (in the order of difficulty to setup): use more client ports by setting net.ipv4.ip_local_port_range to a wider range, use more server ports by asking the web server to listen to several additional ports (81, 82, 83,), use more client IP by configuring additional IP on the load balancer and use them in a round-robin fashion, use more server IP by configuring additional IP on the web server⁶.

But there are ways to tune the TCP/IP stack to reduce the impact of this problem, allowing the system to take advantage of all resources at its disposal. Basically, we can tune several TCP/IP stack parameters, but in this context, there are 2 that really make the difference: The time that the ports is in "waiting" status since it was released from the application, and it's in fact released by the system for reuse. It's called the TIME_WAIT The maximum number of ephemeral ports that the system can use, from the total pool of 65535 available ports. Let's just call it RANGE EPHEMERAL PORTS Reduce the TIME_WAIT by setting the tcp_fin_timeout kernel value on /proc/sys/net/ipv4/tcp_fin_timeout, using the command `echo 30 > /proc/sys/net/ipv4/tcp_fin_timeout` to set it to 30 seconds.

Increase the range of ephemeral ports by setting ip_local_port_range kernel value on /proc/sys/net/ipv4/ip_local_port_range using the command `echo "32768 65535" > /proc/sys/net/ipv4/ip_local_port_range`, this will set the port range from 32768 to 65535.

A common misunderstanding is that a server cannot accept more than 65,536 (2¹⁶) TCP sockets because TCP ports are 16-bit integer numbers. First, the number of ports is limited to 65,536, but this limitation applies only to a single IP address. Supposing that we are limited by the number of ports to have more than 65,536 clients, then adding more IP addresses to the server machine (either by adding new network cards, or simply by using IP aliasing for the existing network card) would solve the problem (even if, for opening 12 million client would need 184 network cards or IP aliases on the server machine). In fact, the misunderstanding comes from the fact that the server does not use its listening IP address and a different ephemeral port for each new socket to distinguish among the sockets, but it uses the same listening IP address and the same listening port for all sockets and it distinguishes among sockets by using the IP address and the ephemeral port of each client. Therefore, MigratoryData Server uses a single port to accept any number of clients and optionally it uses another few ports for JMX monitoring, HTTP monitoring, etc <https://mrotaru.wordpress.com/category/websockets/>

max descriptors per process: Because one cannot increase the maximum number of socket descriptors per process to a value larger than the current kernel maximum (fs.nr_open) and because the kernel maximum defaults to 1048576 (10242), prior to running the ulimit command, we increased the kernel maximum accordingly as follows: `echo 20000500 > /proc/sys/fs/nr_open` CHECK LAB NOTES RECENT TESTS, A LOT <https://mrotaru.wordpress.com/category/websockets/>

The kernel value parameters aren't saved with these commands, and are reset to the default values on system reboot, thus make sure to place the commands on a system startup script such as /etc/rc.local. <https://www.outsystems.com/forums/discussion/6956/how-to-tune-the-tcp-ip-stack-for-high-volume-of-web-requests/>

look for _what can limit concurrent sockets?_ in lab notes

The number of connections wstest can open on a server is limited by the number of ephemeral ports on the machine on the outgoing interface / IP. Something like 64k at most. If you need to test the server with more connections, currently you will need to run multiple instances of wstest (on different machines). from: <http://autobahn.ws/testsuite/usage.html#mode-massconnect>

Understanding tcp sockets: sockets are left in TIME_WAIT on the server when the client suite is killed which makes sense when client exits gracefully no sockets left in TIME_WAIT / CLOSE_WAIT by the server on_close() from tornado only kicks in on client close .onclose in js only kicks in when server initiates through self.close() when server terminates through self.close() sockets ends in TIME_WAIT for a minute (also makes sense according to docs) describe how you can influence that (reusing web sockets + delay time)

ulimit, nofile, fs.maxfile, net.ipv4.ip_local_port_range = 32768 65535 to /etc/sysctl.conf for ip ranges on the clients

4.4 Baseline architecture on a local network

Architecture diagrams. Technical details - local server capabilities. Load testing. LOCAL NETWORK DESCRIPTION

4.5 Cloud architecture setup

selected aws

multiple availability zones very important in a shared cloud environment - often because another customer is getting hit by a DDoS.[GB14]

In Amazons environment, Grozev and Buyya suggest using Elasticache service [GB14]. We cannot use this, also described in 4.5.

4.5.1 Route53

dns mapped to instance ips / load balancer dns names lbr, geo, weighed rr their internal heuristics take network conditions from past weeks [VERIFY], rather ambiguous For each end users location, Route 53 will return the most specific Geo DNS record that includes that location. In other words, for a given end users location, Route 53 will first return a state record; if no state record is found, Route 53 will return a country record; if no country record is found, Route 53 will return a continent record; and finally, if no continent record is found, Route 53 will return the global record.

health checks - can be used custom-defined, but for us we can use load balancing health checks (will only route if healthy instances behind it) [http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/dns-failover-complex-configs.html]

http://docs.aws.amazon.com/Route53/latest/DeveloperGuide/Tutorials.html

4.5.2 Autoscaling

instances behind each groups. automatically scalable, connected to external redis and mysql instances.

only within one region, multiple availability zones http://docs.aws.amazon.com/AutoScaling/latest/DeveloperGuide/G

http://docs.aws.amazon.com/AutoScaling/latest/DeveloperGuide/how-as-works.html#arch-AutoScalingMultiAZ

Auto Scaling attempts to distribute instances evenly between the Availability Zones that are enabled for your Auto Scaling group. Auto Scaling does this by attempting to launch new instances in the Availability Zone with the fewest instances. alarm - object that watches over a single metric (e.g. avg cpu use of ec2 instances in auto scaling group over specified time period) (OK, ALARM, INSUFFICIENT_DATA), can trigger scaling up / down policy on an autoscaling group policy variants (change ExactCapacity, ChangeInCapacity, PercentChangeInCapacity) recommendation one policy for scaling out, another policy for scaling in http://docs.aws.amazon.com/AutoScaling/latest/DeveloperGuide/as-scale-based-on-demand.html http://docs.aws.amazon.com/AmazonCloudWatch/latest/DeveloperGuide/AlarmThatSer

http://docs.aws.amazon.com/AmazonCloudWatch/latest/DeveloperGuide/US_AlarmAtThresholdEC2.html

Alarms

avg / min / max / sum / sample count (for whole group) of cpu / disk io / network io is $i = i / i$ / $i =$ than x for at least X consecutive periods of 5 min / 15 min / 1 hour / 6 hours

Scaling Policies

execute policy when alarm add / remove / set to X instances / and then wait Y (cooldown period)

Launch configurations

4.6 Load balancing

elb uses dns name to server a pool of lb instances in the backend (you might exceed one instance connection limit), that's why they are not provided a static ip routing: request count-based for http(s), for others (tcp which is a workaround for websockets) they use tcp [<http://docs.aws.amazon.com/ElasticLoadBalancing/latest/elb-works.html#request-routing>] The client uses DNS round robin to determine which IP address to use to send the request to the load balancer. No control over routing without load balancing groups. For example, if you have ten instances in Availability Zone us-west-2a and two instances in us-west-2b, the traffic is equally distributed between the two Availability Zones. As a result, the two instances in us-west-2b serve the same amount of traffic as the ten instances in us-west-2a. Instead, you should distribute your instances so that you have six instances in each Availability Zone. To distribute traffic evenly across all back-end instances, regardless of the Availability Zone, enable cross-zone load balancing on your load balancer. However, we still recommend that you maintain approximately equivalent numbers of instances in each Availability Zone for higher fault tolerance. [<http://docs.aws.amazon.com/ElasticLoadBalancing/latest/DeveloperGuide/how-elb-works.html#request-routing>] only 64k connections possible per ELB instance TODO: VERIFY THAT stickiness only works for HTTP/HTTPS protocols Once you have a testing tool in place, you will need to define the growth in the load. We recommend that you increase the load at a rate of no more than 50 percent every five minutes. Both step patterns and linear patterns for load generation should work well with Elastic Load Balancing. If you are going to use a random load generator, then it is important that you set the ceiling for the spikes so that they do not go above the load that Elastic Load Balancing will handle until it scales (see Pre-Warming the ELB). [<https://aws.amazon.com/articles/1636185810492479#pre-warming>]

4.7 Data Layer

Technologies we used described in 3.3. We used Mysql and redis as a caching solution.

4.7.1 Amazon Relational Database Service

Does not support mysql cluster as of July 2015. Support coming but not released yet. Offers mysql read replicas. With a single master. That obviously scales only reads You can create a MySQL Read Replica in a different region than the source DB instance to improve your disaster recovery capabilities, scale read operations into a region closer to end users, or make it easier to migrate from a data center in one region to a data center in another region.

4.7.2 Elasticache

doesn't support crossregion (you cannot connect to an elasticache cluster from outside of the region)

The above explain why, for a scaled out solution we had to roll with our custom setup. We were suggested using technologies like DynamoDB and Kinesis, but out of scope.

4.7.3 Baseline architecture deployed in the cloud

4.7.4 Improved architecture deployed in the cloud

4.8 Experiment deliverables

Chapter 5

Scalability Measurements

Here we describe the metrics we settled for.

Most of what is necessary for the purpose of this research has been covered by Pushkala Pattabiraman et al [MMPROJ3]. They realized that cloud computing and its measurement provides a new set of challenges when it comes to measuring performance testing, as opposed to measuring performance of traditional software systems. They list some key points for measuring cloud applications, among them we can find: validating and ensuring the elasticity of scalability and evaluating utility service billings and pricing models. The latter is also important in my case since cost is one of the driving factors in assessing the scalability of my system. A question they raise regarding this is: How to use a transparent approach to monitor and evaluate the correctness of the utility bill based on a posted price model during system performance evaluation and scalability measurement?. The authors divide the performance indicators into three groups: computing resource (CPU, disk, memory, networks) - they can be helpful in establishing baseline architecture in my case, workload indicators (connected users, throughput and latency), performance indicators - processing speed, system reliability and scalability based on the given QoS standards. For each they propose formal models with pluggable values and graphic representations (BELOW). On top of that, the research contains a case study performed in the Amazon EC2 environment [MMPROJ3]. Cloud limitations need to be taken into account. One needs to be aware of hidden costs (e.g. autoscaling service is free on EC2, but it requires cloudwatch, which is not). The authors also advise to pay attention to inconsistencies in performance and scalability data [MMPROJ3].

many others: There is much more work related to the general scalability of distributed systems. Srinivas and Janakiram in their work [MMPROJ5] mention a metric evaluating scalability as a product of throughput and response time (or any value function) divided by the cost factor. They propose another model considering scalability as a function of synchronization, consistency, availability, workload and faultload. It aims on identifying bottlenecks and hence improving the scalability. The authors also emphasize the fact of interconnectedness of synchronization, consistency and availability. Jogalekar and Woodside [6] propose a strategy-based scalability metric based on cost effectiveness (a function of system's throughput and its quality of service). It separates evaluation of throughput or quantity of work from QoS (which, according to the authors, can be any suitable expression).[MMPROJ6]

PASA[MMPROJ17]

5.1 Selected metrics

latency (messaging + handhsaking) (our suite) sustainable concurrent websocket connectoins (our suite) infrastructure cost (per unit of time, supporting the same number of users) (manual calculation) dropped connections (our suite) cpu usage (cloudwatch) memory usage (our custom server metrics, pidstat) network in network out (cloudwatch) iops reads + writes (bots number and byte values, cloudwatch)

to calculate manually: percentage of available ports used? how long to run on a x euros

availability (SLAs???????) architecture reaction speed to changes in demand (scaling up and scaling

down) (???)

CRAM METRICS TODO: review

5.1.1 EC2 Available metrics

collectible every minute (in detailed mode), by default every 5 minute, paid additionally - delay up to 2 minutes the current EC2 cloud technology does not provide any CloudWatch API to monitor the allocation and utilization of memory for EC2 instances (we do it on our own) limits exist (10 metric, 10 alarms, 1,000,000 million requests, 1000 SNS email notifications per month for free no limits on custom metrics up to 5gb of incoming data logs for free up to 5gb of data archiving for free) but we don't expect to hit them. small delay (up to 2 minutes, that's what we experienced - delays in autoscaling)

alarms can be configured based on this - <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-cloudwatch.html> data not aggregated across regions [lab2]

Instance Specific

[Lab2 - http://docs.aws.amazon.com/AmazonCloudWatch/latest/DeveloperGuide/cloudwatch_concepts.html] custom possible available statistics for each: min, max, sum, avg, sampleCount (count of data points but research that) cpu utilization network in / out (in bytes) disk read / disk write (number of ops + bytes) status checks

Cloud (group) Specific

Collected from Elastic Load Balancer for all instances connected.

Custom Metric Collection

pidstat latencies by our custom client suite for network we looked at ntop, nethogs, nload, vnstat, wireshark but amazon provides this data.

5.2 Selected model

[MMPROJ3] With CRUMs, SPMs etc. four different types of needs: resource utilization and allocation measurement performance measurement under allocated computing resources scalability measurement in different cloud infrastructures cost driven evaluation based on a pre-defined price model

additional values we want to capture to proposed in the model: cost time to scale up / down What we will use to compare architectures: SCM (allocated resources) or SEC (used) (1/cost) as an additional metric to SEC? or amount of s you can run on a 100bucks

Results of all setups get plugged into SCM / SEC and then ESS.

5.3 Baseline, Local network

Lab notes from 11 of May go here:

Basic metrics

- cpu, memory, disk usage (pidstat / CloudWatch)
- network i/o (wireshark, list others analyzed)
- latency, throughput, concurrent connections, messages dropped?

5.3.1 Load Testing

<http://stackoverflow.com/questions/3871886/ssl-and-load-balancing>

elb uses dns name to server a pool of lb instances in the backend (you might exceed one instance connection limit) in order for the full range of ELB machine IP addresses to be utilized, make sure each [client] refreshes their DNS resolution results every few minutes. One test client equals one public IP address. ELB machines seem to route all traffic from a single IP address to the same back-end instance, so if you run more than one test client process behind a single public IP address, ELB regards these as a single client. Use 12 test clients for every availability zone you have enabled on the ELB. These test clients do not need to be in different availability zones they do not even need to be in EC2 (although it is quite attractive to use EC2 instances for test clients). If you have configured your ELB to balance among two availability zones then you should use 24 test clients. Each test client should gradually ramp up its load over the course of a few hours. Each client can begin at (for example) one connection per second, and increase its rate of connections-per-second every X minutes until the load reaches your desired target after a few hours. - WE CANNOT DO THAT

Tools described here [?]

Users Distribution

The authors suggest choosing randomly when generating load as to which operation to perform, on what data size etc. They suggest using different random distributions: uniform, zipfian, latest, multinomial[CSE⁺10]. Another distribution is suggested by Grozev and Buyya [GB14] - Poisson distribution with a constant mean.

5.4 Baseline, Deployed in the cloud

5.5 Improved, Deployed in the cloud

Chapter 6

Experiment results

Chapter 7

Evaluation

7.1 Statistical Significance

<http://technology.stitchfix.com/blog/2015/05/26/significant-sample/>

7.2 Threats to validity

ec2 whacky, control over hardware released to amazon (can be dedicated machines) shared environment, someone else might get ddosed or sth - you have no control over that network is inherently non-deterministic, load tests needed to actually test availability

Chapter 8

Conclusion

Chapter 9

Further work

Chapter 10

Everything below that needs to be removed (except for bib)

Chapter 11

Front Matter

The first thing is to connect the class by saying:

```
\documentclass{uvamscse}
```

11.1 Title

Specify the title of the thesis with `\title` and `\subtitle` commands:

```
\title{MetaThesis}
\subtitle{A Thesis Template Leading by Example}
```

Any thesis can survive without a `\subtitle`, but the `\title` is mandatory.

11.2 Author

Introduce yourself with `\author` and `\authemail`:

```
\author{Vadim Zaytsev}
\authemail{vadim@grammarware.net}
```

Again, `\authemail` is not mandatory. If you need anything fancier, just put it inside `\author`.

```
\author{Vadim Zaytsev\footnote{Yes, that one.}}
```

The footnote would be printed on the bottom of the title page, and will be referred to by a symbol, not by a number as any footnotes within the main document body.

11.3 Date

By default, the date inserted in your PDF is the day of the build, e.g., “March 25, 2014”. If you want it to be formatted differently or be more vague or outright fake, use `\date`:

```
\date{Spring 2014}
```

The argument is just a string, the format is unrestricted:

```
\date{Tomorrow. Honestly.}
```

11.4 Host

If your hosting organisation is not the UvA, specify it with `\host`. The logo on the bottom of the title page will still be the UvA one, because this is the organisation guaranteeing your degree.

```
\host{Grammarware, Inc., \url{http://grammarware.github.io}}
```

NB: footnotes will not work, unless you know how to `\protect` them.



Figure 11.1: A hypothetical thesis title page without a cover picture (on the left), with an overly large one (in the centre) and with a tiny pic (on the right).

11.5 Cover picture

If the first page of your thesis looks too blunt, add a picture to it:

```
\coverpic{figures/terminal.png}
```

You can even specify the picture's width as an optional argument:

```
\coverpic[100pt]{figures/terminal.png}
```

How these three options look, you can see from [Figure 11.1](#).

11.6 Abstract

A thesis is fine without an abstract, if you do not feel like writing it and your supervisor does not feel like enforcing it. If you do want an abstract, make it with the `\abstract` command:

```
\abstract{This is not a thesis.}
```

The abstract is just like any other section of your thesis, so you can use any \LaTeX tricks there. If you think that the name “abstract” is too abstract for your abstract, you can still use `\abstract` without being too abstract:

```
\abstract[Confession]{I am a cenosillicaphobiac.}
```

Kent Beck [[JBB⁺93](#)] proposes to have four sentences in a good abstract:

1. The first states the problem.
2. The second states why the problem is a problem.
3. The third is the startling sentence.
4. The fourth states the implication of the startling sentence.

In practice, each of these “sentences” can be longer than an actual sentence, but it is in general a good rule of thumb to condense the summary of your thesis into these four tiny messages. Do not write too much, make it tweetable.

Chapter 12

Core Chapters

The structure of your thesis is up to you and your supervisor. Whatever you do, do not consider the guidelines below as dogmas.

12.1 Classic structure

Problem statement and motivation. You describe in detail what problem the research is addressing, and what is the motivation to address this problem. There is a concise and objective statement of the research questions, hypotheses and goals. It is made clear why these questions and goals are important and relevant to the world outside the university (assuming it exists). You can already split the main research question into subquestions in this chapter. This section also describes an analysis of the problem: where does it occur and how, how often, and what are the consequences? An important part is also to scope the research: what aspects are included and what aspects are deliberately left out, and why?

Research method. Here you describe the methods used to answer the research questions. A good structure of this section often follows the subquestions by providing a method for each. The research method needs a thorough motivation grounded in theory in order to be acceptable. As a part of the method, you can introduce a number of hypotheses — these will be tested by the research, using the methods described here. An important part of this section is validation. How will you evaluate and validate the outcomes of the research?

Background and context. This chapter contains all the information needed to put the thesis into context. It is common to use a revised version of your literature survey for this purpose. It is important to explicitly refer from your text to sources you have used, they will be listed in your bibliography. For example, you can write “A small number of programming languages account for most language use [MR13]”, where the following entry would be included in your bibliography:

[MR13] Leo A. Meyerovich and Ariel S. Rabkin. Empirical Analysis of Programming Language Adoption. In *Proceedings of the 2013 ACM SIGPLAN International Conference on Object Oriented Programming Systems Languages and Applications*, OOPSLA, pages 1–18. ACM, 2013. doi:10.1145/2509136.2509515.

Have a look at § 13 to learn more about citation.

Research. This chapter reports on the execution of the research method as described in an earlier chapter. If the research has been divided into phases, they are introduced, reported on and concluded individually. If needed, this chapter could be split up to balance out the sizes of all chapters.

Results. This chapter presents and clarifies the results obtained during the research. The focus should be on the factual results, not the interpretation or discussion. Tables and graphics should be used to increase the clarity of the results where applicable.

Analysis and conclusions. This chapter contains the analysis and interpretation of the results. The research questions are answered as best as possible with the results that were obtained. The analysis also discussed parts of the questions that were left unanswered.

An important topic is the validity of the results. What methods of validation were used? Could the results be generalised to other cases? What threats to validity can be identified? There is room here to discuss the results of related scientific literature here as well. How do the results obtained here relate to other work, and what consequences are there? Did your approach work better or worse? Did you learn anything new compared to the already existing body of knowledge? Finally, what could you say in hindsight on the research approach by followed? What could have done better? What lessons have been learned? What could other researchers use from your experience? A separate section should be devoted to “future work”, i.e., possible extension points of your work that you have identified. Even other researchers should be able to use those as a starting point.

12.2 Reporting on replications

Here are the guidelines to report on replicated studies [Car10]:

Information about the original study

Research question(s) that were the basis for the design

Participants, their number and any other relevant characteristics

Design as a graphical or textual description of the experimental design

Artefacts, the description of them and/or links to the artefacts used

Context variables as any important details that affected the design of the study or interpretation of the results

Summary of the results in a brief overview of the major findings

Information about the replication

Motivation for conducting the replication as a description of why the replication was conducted: to validate the results, to broaden the results by changing the participant pool or the artifacts.

Level of interaction with original experimenters. The level of interaction between the original experimenters and the replicators should be reported. This interaction could range from none (i.e. simply read the paper) to them being the same people. There is quite a lot of discussion of the level of interaction allowed for the replication to be “successful”, but this level should be reported even without addressing the controversy.

Changes to the original experiment. Any changes made to the design, participants, artifacts, procedures, data collected and/or analysis techniques should be discussed along with the motivation for the change.

Comparison of results to original

Consistent results, when replication results supported results from the original study, and

Differences in results, when results from the replication did not coincide with the results from the original study. Authors should also discuss how changes made to the experimental design (see above) may have caused these differences.

Drawing conclusions across studies

NB: this section contains portions of text repeated directly from Carver [Car10] and only slightly massaged. Do not do this for your thesis, write your own thoughts down.

12.3 L^AT_EX details

12.3.1 Environments

A L^AT_EX environment is something with opening and closing tags, which look like `\begin{name}` and `\end{name}`. Some useful environments to know:

<code>itemize</code>	bullet lists
<code>enumerate</code>	numbered lists
<code>description</code>	definition lists
<code>center</code>	centered line elements
<code>flushright</code>	right aligned lines
<code>flushleft</code>	left aligned lines
<code>tabular</code>	table
<code>longtable</code>	multi-page table (needs the <code>longtable</code> package)
<code>sideways</code>	rotates some text
<code>quote</code>	block quote
<code>verbatim</code>	unformatted text
<code>minipage</code>	compound box with elements inside
<code>boxedminipage</code>	compound box with elements inside and a border around it
<code>table</code>	floating table (needs to have <code>tabular</code> nested inside)
<code>figure</code>	floating figure
<code>sourcecode</code>	floating listing
<code>equation</code>	mathematical equation
<code>lstlisting</code>	pretty-printed syntax highlighted listing
<code>multline</code>	mathematical equation spanning over multiple lines
<code>eqnarray</code>	system of mathematical equations
<code>gather</code>	bundled mathematical equations
<code>align</code>	bundled and aligned mathematical equations
<code>array</code>	matrix
<code>CD</code>	commutative diagrams

12.4 Listings

```
1 define(Ps1,G1,G2)
2   ←
3   usedNs(G1,Uses),
4   ps2n(Ps1,N),
5   require(
6     member(N,Uses),
7     'Nonterminal ~q must not be fresh.',
8     [N]),
9   new(Ps1,N,G1,G2),
10  !.
```

Listing 12.1: Code in Prolog

```

module Syntax

imports Numbers
imports basic/Whitespace

exports
sorts
  Program Function Expr Ops Name Newline

context-free syntax
  Function+ → Program
  Name Name+ "=" Expr Newline+ → Function
  Expr Ops Expr → Expr {left,prefer,cons(binary)}
  Name Expr+ → Expr {avoid,cons(apply)}
  "if" Expr "then" Expr "else" Expr → Expr {cons(ifThenElse)}
  "(" Expr ")" → Expr {bracket}
  Name → Expr {cons(argument)}
  Int → Expr {cons(literal)}
  "_" → Ops {cons(minus)}
  "+" → Ops {cons(plus)}
  "==" → Ops {cons(equal)}

```

Listing 12.2: Code in SDF

```

1 import types.*;
2 import org.antlr.runtime.*;
3
4 public class TestEvaluator
5   public static void main(String[] args) throws Exception {
6
7     // Parse file to program
8     ANTLRFileStream input = new ANTLRFileStream(args[0]);
9     FLLexer lexer = new FLLexer(input);
10    CommonTokenStream tokens = new CommonTokenStream(lexer);
11    FLParser parser = new FLParser(tokens);
12    Program program = parser.program();
13
14    // Parse sample expression
15    input = new ANTLRFileStream(args[1]);
16    lexer = new FLLexer(input);
17    tokens = new CommonTokenStream(lexer);
18    parser = new FLParser(tokens);
19    Expr expr = parser.expr();
20
21    // Evaluate program
22    Evaluator eval = new Evaluator(program);
23    int expected = Integer.parseInt(args [2]);

```

Listing 12.3: Code in Java

```

1  #!/usr/local/bin/python
2  # wiki: BGF
3  import os
4  import sys
5  import slpsns
6  import elementtree.ElementTree as ET
7
8  # root::nonterminal* production*
9  class Grammar:
10     def __init__(self):
11         self.roots = []
12         self.prods = []
13     def parse(self, fname):
14         self.roots = []
15         self.prods = []
16         self.xml = ET.parse(fname)
17         for e in self.xml.findall('root'):
18             self.roots.append(e.text)
19         for e in self.xml.findall(slpsns.bgf_('production')):
20             prod = Production()
21             prod.parse(e)
22             self.prods.append(prod)

```

Listing 12.4: Code in Python

Chapter 13

Literature

BIBTeX is a JSON-like format for bibliographic entries. Encode each source once as a BIBTeX entry, give it a name and refer to it from any place in your thesis. The bibliography at the end of the thesis will be compiled automatically from those entries that are referenced at least once, it will also be automatically sorted and fancyfied (URLs, DOIs, etc).

DOI is a digital object identifier, it is uniquely and immutably assigned to any paper published in a well-established journal or conference proceedings and can be used to refer to it. When used in a browser, it resolves to a publisher's website where paper can be obtained. Including DOIs in citations is considered good practice and lets the readers of your thesis get to the text of the paper in one click. Books do not have DOIs, only ISBNs; some workshop proceedings and most unofficial publications do not have DOIs. If you want to get a DOI assigned to your work such as a piece of code, upload it to [FigShare](#).

Keys in key-value pairs within each BIBTeX entry are never quoted, values usually are, but can also be included within curly brackets or left as is, which works fine for numbers (e.g., years). If you want to preserve the value from any adjustments (e.g., no recapitalisation in titles), use curly braces *and* quotes. Separate authors and editors by “and”, which will automatically be mapped to commas or left as “and”s as necessary.

13.1 Books

[GJ08] is just as good as the Dragon Book, but newer and has an awesome extended bibliography available for free.

```
@book{GruneJacobs,
  author   = "D. Grune and C. J. H. Jacobs",
  title    = "{Parsing Techniques: A Practical Guide}",
  series   = "Monographs in Computer Science",
  edition  = 2,
  publisher = "Springer",
  url      = "http://www.cs.vu.nl/~dick/PT2Ed.html",
  year     = 2008,
}
```

13.2 Journal papers

Not all TOSEM papers are hard to read [KLV05].

```
@article{GrammarwareAgenda,
  author    = "Paul Klint and Ralf L{\a}mmel and Chris Verhoef",
  title     = "{Toward an Engineering Discipline for Grammarware}",
  journal   = "ACM Transactions on Software Engineering Methodology (TOSEM)",
  volume    = 14,
  number    = 3,
  year      = 2005,
  pages     = "331--380",
}
```

13.3 Conference papers

There is no limit to how many grammars can be used in one paper, but the current record stands at 569 [Zay13].

```
@inproceedings{Micropatterns2013,
  author = "Vadim Zaytsev",
  title = "{Micropatterns in Grammars}",
  booktitle = "{Proceedings of the Sixth International Conference on Software Language Engineering (SLE 2013)}",
  year = 2013,
  editor = "Martin Erwig and Richard F. Paige and Eric Van Wyk",
  volume = "8225",
  series = "LNCS",
  pages = "117--136",
  address = "Switzerland",
  month = oct,
  publisher = "Springer International Publishing",
  doi = "10.1007/978-3-319-02654-1_7",
}
```

13.4 Theses

The seventh PhD student of Paul Klint was Jan Rekers [Rek92].

```
@phdthesis{Rekers92,
  author = "J. Rekers",
  title = "{Parser Generation for Interactive Environments}",
  school = "University of Amsterdam",
  year = 1992,
  url = "http://homepages.cwi.nl/~paulk/dissertations/Rekers.pdf",
}
```

There is also `mastersthesis` type with exactly the same structure for referring to Master's theses.

13.5 Technical reports

The original seminal work introducing two-level grammars was never published in any book or conference, but there is a technical report explaining it [vW65]. SMC, or *Stichting Mathematisch Centrum*, was the old name of CWI fifty years ago.

```
@techreport{Wijngaarden65,
  author = "Adriaan van Wijngaarden",
  title = "{Orthogonal Design and Description of a Formal Language}",
  month = oct,
  year = 1965,
  institution = "SMC",
  type = "{MR 76}",
  url = "http://www.fh-jena.de/~kleine/history/languages/VanWijngaarden-MR76.pdf",
}
```

13.6 Wikipedia

You do not refer to Wikipedia from academic writing, it works the other way around.

13.7 Anything else

You can refer to pretty much anything (websites, blog posts, software) through `misc` type of entry [\[Par08\]](#):

```
@misc{ANTLR,
  author      = "Terence Parr",
  title       = "{ANTLR}---ANother Tool for Language Recognition",
  howpublished = "Software",
  url         = "http://antlr.org",
  year        = "2008"
}
```


Bibliography

- [Aga11] S. Agarwal. Toward a push-scalable global internet. In *Computer Communications Workshops (INFOCOM WKSHPS), 2011 IEEE Conference on*, pages 786–791. IEEE, 2011. doi:[10.1109/INFCOMW.2011.5928918](https://doi.org/10.1109/INFCOMW.2011.5928918).
- [Amz02] C. Amza. Specification and implementation of dynamic Web site benchmarks. In *Workload Characterization, 2002. WWC-5. 2002 IEEE International Workshop on*, pages 3–13. IEEE, November 2002. doi:[10.1109/WWC.2002.1226489](https://doi.org/10.1109/WWC.2002.1226489).
- [Car10] Jeffrey C. Carver. Towards Reporting Guidelines for Experimental Replications: A Proposal. In *Proceedings of the International Workshop on Replication in Empirical Software Engineering Research*, RESER, May 2010.
- [CL12] Oscar Cassetti and Saturnino Luz. The websocket api as supporting technology for distributed and agent-driven data mining oscar cassetti, 2012.
- [CSE⁺10] Brian F. Cooper, A. Sibley, E. Tam, R. Ramakrishnan, and R. Sears. Benchmarking cloud serving systems with YCSB. In *SoCC '10 Proceedings of the 1st ACM symposium on Cloud computing*, pages 143–154. ACM, 2010. doi:[10.1145/1807128.1807152](https://doi.org/10.1145/1807128.1807152).
- [GB14] Nikolay Grozev and Rajkumar Buyya. Multi-Cloud Provisioning and Load Distribution for Three-Tier Applications. *ACM Transactions on Autonomous and Adaptive Systems (TAAS)*, 9(13), October 2014. URL: <http://dl.acm.org/citation.cfm?id=2662112>.
- [GJ08] Dick Grune and C. J. H. Jacobs. *Parsing Techniques: A Practical Guide*. Monographs in Computer Science. Springer, 2 edition, 2008. URL: <http://www.cs.vu.nl/~dick/PT2Ed.html>.
- [GL02] Seth Gilbert and Nancy Lynch. Brewer’s conjecture and the feasibility of consistent, available, partition-tolerant web services. *ACM SIGACT News*, 33(2):51–59, June 2002. URL: <http://dl.acm.org/citation.cfm?id=564601>.
- [Hei10] Gernot Heiser. Systems benchmarking crimes. Toolkit website, January 2010. URL: <https://www.cse.unsw.edu.au/~gernot/benchmarking-crimes.html>.
- [How15] Ed Howorka. Colocation beats the speed of light. Website, February 2015. URL: <http://edhoworka.com/colocation-beats-the-sol/>.
- [HR83] Theo Harder and Andreas Reuter. Principles of transaction-oriented database recovery. *ACM Computing Surveys (CSUR)*, 15(4):287–317, December 1983. doi:[10.1145/289.291](https://doi.org/10.1145/289.291).
- [JBB⁺93] Ralph E. Johnson, Kent Beck, Grady Booch, William R. Cook, Richard P. Gabriel, and Rebecca Wirfs-Brock. How to Get a Paper Accepted at OOPSLA. In Timlynn Babitsky and Jim Salmons, editors, *Proceedings of the Eighth Annual Conference on Object-Oriented Programming Systems, Languages and Applications*, OOPSLA, pages 429–436. ACM, 1993.
- [KLV05] Paul Klint, Ralf Lämmel, and Chris Verhoef. Toward an Engineering Discipline for Grammarware. *ACM Transactions on Software Engineering Methodology (TOSEM)*, 14(3):331–380, 2005.

- [LD12] Louis Liu and Alexander Dekhtyar. CSC 560: Advanced Topics in Databases: Modern DBMS Architectures, Eventual Consistency. University website, December 2012. URL: https://wiki.csc.calpoly.edu/csc560/raw.../wiki/.../Louis_Liu_paper.pdf.
- [Lei08] Tom Leighton. Improving Performance on the Internet. *ACM Queue - Scalable Web Services*, 6(6):20–29, October 2008. URL: <http://queue.acm.org/detail.cfm?id=1466449>.
- [MR13] Leo A. Meyerovich and Ariel S. Rabkin. Empirical Analysis of Programming Language Adoption. In *Proceedings of the 2013 ACM SIGPLAN International Conference on Object Oriented Programming Systems Languages and Applications*, OOPSLA, pages 1–18. ACM, 2013. doi:10.1145/2509136.2509515.
- [Ora15] Oracle. Guide to Scaling Web Databases with MySQL Cluster. White Paper, April 2015. URL: <https://www.mysql.com/why-mysql/white-papers/guide-to-scaling-web-databases-with-mysql-cluster/>.
- [Par08] Terence Parr. ANTLR—ANother Tool for Language Recognition. Toolkit website, 2008. URL: <http://antlr.org>.
- [Qve10] Nikolai Qveflander. *Pushing real time data using HTML5 Web Sockets*. PhD thesis, Ume University, 2010. URL: <http://citeseerx.ist.psu.edu/viewdoc/summary?doi=10.1.1.175.2195>.
- [RBR11] Nicolas Ruffin, Helmar Burkhart, and Sven Rizotti. Social-data storage-systems. In *DBSozial '11 Databases and Social Networks*, pages 7–12. ACM, 2011. doi:10.1145/1996413.1996415.
- [Rek92] J. Rekers. *Parser Generation for Interactive Environments*. PhD thesis, University of Amsterdam, 1992. URL: <http://homepages.cwi.nl/~paulk/dissertations/Rekers.pdf>.
- [SSS⁺08] Will Sobel, Shanti Subramanyam, Akara Sucharitakul, Jimmy Nguyen, Hubert Wong, Arthur Klepchukov, Sheetal Patil, O. Fox, and David Patterson. Cloudstone: Multi-platform, multi-language benchmark and measurement tools for web 2.0, 2008.
- [vW65] Adriaan van Wijngaarden. Orthogonal Design and Description of a Formal Language. MR 76, SMC, October 1965. URL: <http://www.fh-jena.de/~kleine/history/languages/VanWijngaarden-MR76.pdf>.
- [WG06] Charles Weinstock and John Goodenough. On system scalability. Technical Report CMU/SEI-2006-TN-012, Software Engineering Institute, Carnegie Mellon University, Pittsburgh, PA, 2006. URL: <http://resources.sei.cmu.edu/library/asset-view.cfm?AssetID=7887>.
- [WS04] Lloyd G. Williams and Connie U. Smith. Web Application Scalability: A Model-Based Approach. *Software Engineering Research and Performance Engineering Services*, 2004. URL: <http://www.perfeng.com/papers/scale04.pdf>.
- [Zay13] Vadim Zaytsev. Micropatterns in Grammars. In Martin Erwig, Richard F. Paige, and Eric Van Wyk, editors, *Proceedings of the Sixth International Conference on Software Language Engineering (SLE 2013)*, volume 8225 of *LNCS*, pages 117–136, Switzerland, October 2013. Springer International Publishing. doi:10.1007/978-3-319-02654-1_7.